

Better AI through Logical Scaffolding

Nikos Aréchiga¹, Jonathan DeCastro¹, Soonho Kong¹, and Karen Leung²

¹ Toyota Research Institute, Los Altos, CA and Cambridge, MA, USA

² Stanford University, Palo Alto, CA, USA

Abstract. We describe the concept of *logical scaffolds*, which can be used to improve the quality of software that relies on AI components. We explain how some of the existing ideas on runtime monitors for perception systems can be seen as a specific instance of logical scaffolds. Furthermore, we describe how logical scaffolds may be useful for improving AI programs beyond perception systems, to include general prediction systems and agent behavior models.

Keywords: AI · Autonomous systems · Formal methods.

1 Introduction

Recent progress in AI has led to possible deployment in a wide variety of important domains. This includes safety-critical cyberphysical systems such as automobiles [1] and airplanes [7], but also decision making systems in diverse domains including legal [15] and military applications [3].

Current AI programs differ from traditional programs in their reliance on data. The specification, input-output semantics, and executable generation procedure are all data driven [9].

Unlike traditional software development, in AI programs a specification is not formally articulated. Indeed, in many of the most promising recent applications of AI, such as vision and human intent prediction, it is not feasible to write a formal specification. Instead, an implicit specification is provided via a test set, and the goal is to achieve a certain performance over the test set.

Traditional software development specifies the input-output semantics of the program in a programming language. In AI programs, the engineer provides a training dataset, and the program must match the input-output statistics of the dataset.

Instead of using a compiler to translate the programming language constructs to machine code, an engineer provides a “skeleton” in the form of a powerful function approximator (such as a neural network). The engineer then uses an optimization procedure to search for the parameterization that provides the best approximation to the input-output statistics of the training data. A cross-validation set is used to check generalizability of the learned function to unseen data.

This paradigm has proven to be powerful, especially in domains in which it is difficult to formally articulate the specification for the program, much less

write a declarative program. However, this approach suffers from the following drawbacks.

1. **Implicit specifications** Since the specification is given implicitly as a desired performance over a test set, it is difficult for the tests to ascertain whether the program is providing the right answers for the right reasons [10]. For this reason, the test set may fail to test the right things. This type of implicit specification is inadequate for use in a safety case, and it will be difficult to ascertain that programs tested in this way will be safe to deploy.
2. **Non-representative training set** Since the program seeks to match the statistical input-output properties of the training dataset, deficiencies of this training set will extend to deficiencies of the program. For example, scenarios that occur rarely in the training set may be fairly common in the real world, leading to degraded performance in deployment [2]. In this sense, the training set may fail to train for realistic scenarios.
3. **Robustness and sensitivity to adversarial attacks** Since the model is templated by a functional template with many degrees of freedom, it is common for the process to result in programs that are susceptible to extreme sensitivity to irrelevant features of the input. The literature on *adversarial examples* demonstrates how slight perturbations to an input can lead to incorrect results with high confidence [6].

We propose to attack these issues by the use of *logical scaffolds*, which are lightweight formal properties that provide some information about the relationship of the program inputs and outputs. These logical scaffolds can be written in languages for which monitoring algorithms exist, such as Signal Temporal Logic [12], Signal Convolutional Logic [16], Timed Quality Temporal Logic [4] and many others. Logical scaffolds may arise from a number of different sources, including a formalization of physical laws, domain knowledge, and common sense.

Logical scaffolds are more general than related work such as reasonableness monitors [5] and model assertions [8] because scaffolds can be used for different types of AI programs beyond merely perception. Furthermore, recent work in smoothly differentiable formulations of STL and MTL [11, 14, 13] enable the scaffolds to become part of the training process directly.

2 Logical scaffolds

Informally, a logical scaffold is a predicate that encodes something that is believed to be true about the input-output relation of an AI program. It is *not* a complete specification. If a specification existed, the scaffold would be a logical consequence of the complete specification. In other words, the scaffold is a consequence of correct functionality. As such, it constitutes a necessary, but not sufficient, condition for correct behavior.

Most of the existing literature on monitoring runtime properties is centered around perception systems. These runtime monitors can be formalized as scaffolds, but the key idea can be generalized beyond perception to include applications in explainable intent prediction and expressible behavior modeling.

Sources of logical scaffolds are as diverse as the sources of human intuition about the application domain. The following is not an exhaustive list.

– **Perception**

- Commonsense notions of label consistency, like the properties monitored in [4] and [8], in which class labels are not expected to mutate or drop out between frames.
- Class-specific information, such as the intuition that a mailbox should not be seen crossing the road [5].

– **Intent prediction and behavior modeling**

- Physics-derived knowledge, such as knowledge of maximum speeds or actuator capabilities, for example, that on an icy road, other vehicles may be out of control or less able to brake and swerve.
- Natural expectations that pedestrians and vehicles are unlikely to seek damage to themselves, unless they are out of control.

The challenges related to implicit specifications are ameliorated by using logical scaffolds at training and testing time. For generative models, such as intent predictors and agent models, we can impose an understandable structure on the latent space, as described below. At testing time, we are able to use parametric logical scaffolds to learn explanations of the input-output behavior of the system, which can be used to check that the system is passing its tests for the right reasons.

The challenges related to non-representative training sets as well as robustness and sensitivity to adversarial attacks can be ameliorated by using the scaffolds at deployment.

The work of Kang et al. [8] has shown how hand crafted runtime monitors can be used to flag scenarios in which the program fails. These scenarios can then be added to the training set, yielding greatly improved performance. This approach has a flavor of active learning, in which the scenarios that are most difficult for the program are fed back for further study. Here, we propose that the monitors need not be hand crafted, but automatically generated from logical scaffolds that express a variety of properties. Conversely, in [4], Dokhanchi et al. automatically generate monitors from Timed Quality Temporal Logic that check for label stability, i.e., ensuring that labels do not mutate across frames.

2.1 Training

There may be many ways that logical scaffolds can be used at training time. In this work we consider training generative models that make use of a latent space.

Generative models are models that can generate data that is similar to the data they are trained on. Important examples of generative models for contemporary applications include the following.

1. Intent predictors, which are used by autonomous vehicles or other robots to predict future trajectories of other agents, such as automobiles, pedestrians and bicycles.
2. Reactive agents, which seek to generate appropriate behaviors for a specific environment. Examples may include simulations agents that subject an autonomous vehicle, drone, or robot to challenging but realistic behaviors, as well as decision-making agents for different applications.
3. Scenario generators, which seek to synthesize testing and simulation environments that may be challenging but still realistic.

We can use logical scaffolds expressed in a differentiable logic (e.g. [11], [13], [14]) and use them to add structure to the latent space. A diagram of this idea is shown in Figure 1.

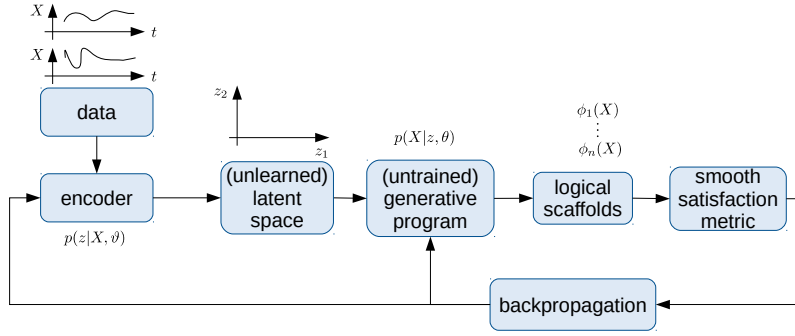


Fig. 1. Learning latent spaces from explanations

2.2 Testing

At test time, we are interested in finding out whether the model is producing the correct outputs for the correct reasons. To accomplish this, we make use of *parametric* logical scaffolds, i.e. scaffolds with free parameters.

We sample from the latent space, and prompt the model to generate an output. Then, we take a bank of pSTL formulas, and fit values for each of the parameters. Then, we check to see which of the parameters have clusters that correspond to the clusters of the original latent space. The corresponding STL formulas are “explanations” of the latent space clusters. A diagram of this idea is shown in Figure 2.

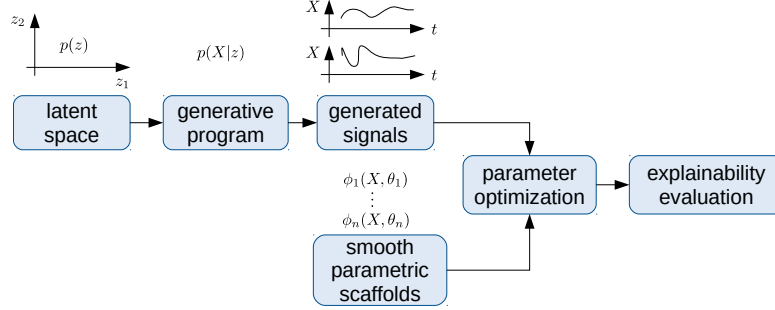


Fig. 2. Learning explanations for latent spaces

2.3 Deployment

At deployment, the logical scaffold can be used to detect anomalies. These anomalies can later be fed back into the training procedure for the next iteration of the system. The use of logical scaffolds for runtime improvement is predicated on the fact that these scaffolds are written in formal languages that support runtime monitoring.

The work of [8] has already demonstrated how hand-crafted runtime monitors can be used to greatly improve the performance of single shot detectors, and the work of [4] has developed a special-purpose runtime logic to monitor the stability of class labels.

The idea of augmenting AI programs with knowledge of physics (“Newton + Hinton”³) is not new. We believe that even greater impact can be obtained from the broader principle of systematically developing logical scaffolds that encode physics domain-specific knowledge, and common sense.

3 Conclusions and Future Work

We have outlined a technique of “logical scaffolding”, which involves conditions that are necessary for correctness, but not sufficient. We have outlined how these logical scaffolds can be used to improve the performance and reliability of AI systems at training, testing, and deployment.

In future work, we will explore case studies in perception, behavior modeling, and scenario generation.

References

1. Mariusz Bojarski, Davide Del Testa, Daniel Dworakowski, Bernhard Firner, Beat Flepp, Praseoon Goyal, Lawrence D. Jackel, Mathew Monfort, Urs Muller, Jiakai

³ Thanks to Adrien Gaidon for this term

- Zhang, Xin Zhang, Jake Zhao, and Karol Zieba. End to End Learning for Self-Driving Cars. *arXiv:1604.07316 [cs]*, April 2016. arXiv: 1604.07316.
2. Yin Cui, Menglin Jia, Tsung-Yi Lin, Yang Song, and Serge Belongie. Class-Balanced Loss Based on Effective Number of Samples. *arXiv:1901.05555 [cs]*, January 2019. arXiv: 1901.05555.
3. M L Cummings. Artificial Intelligence and the Future of Warfare. Technical report, Chatham House, 2017.
4. Adel Dokhanchi, Heni Ben Amor, Jyotirmoy V. Deshmukh, and Georgios Fainekos. Evaluating Perception Systems for Autonomous Vehicles Using Quality Temporal Logic. In *Runtime Verification*, volume 11237. Springer International Publishing, Cham, 2018.
5. Leilani H. Gilpin. Reasonableness Monitors. *AAAI*, 2018.
6. Ian J. Goodfellow, Jonathon Shlens, and Christian Szegedy. Explaining and Harnessing Adversarial Examples. In *arXiv:1412.6572 [cs, stat]*, 2015. arXiv: 1412.6572.
7. Kyle D. Julian, Jessica Lopez, Jeffrey S. Brush, Michael P. Owen, and Mykel J. Kochenderfer. Policy compression for aircraft collision avoidance systems. In *2016 IEEE/AIAA 35th Digital Avionics Systems Conference (DASC)*. IEEE, September 2016.
8. Daniel Kang, Deepti Raghavan, Peter Bailis, and Matei Zaharia. Model Assertions for Debugging Machine Learning. *NeurIPS MLSys Workshop*, 2018.
9. Andrej Karpathy. Software 2.0, November 2017.
10. Philip Koopman and Michael Wagner. Toward a Framework for Highly Automated Vehicle Safety Validation. In *SAE World Congress*, April 2018.
11. Karen Leung, Nikos Arechiga, and Marco Pavone. Backpropagation for Parametric STL. In *Intelligent Vehicles Symposium*, 2019.
12. Oded Maler and Dejan Nickovic. Monitoring Temporal Properties of Continuous Signals. In *Formal Techniques, Modelling and Analysis of Timed and Fault-Tolerant Systems*, 2004.
13. Noushin Mehdipour, Cristian Ioan Vasile, and Calin Belta. Arithmetic-geometric mean robustness for control from signal temporal logic specifications. *CoRR*, abs/1903.05186, 2019.
14. Yash Vardhan Pant, Houssam Abbas, and Rahul Mangharam. Smooth operator: Control using the smooth robustness of temporal logic. In *2017 IEEE Conference on Control Technology and Applications (CCTA)*, pages 1235–1240, Mauna Lani Resort, HI, USA, August 2017. IEEE.
15. Walter L. Perry, Brian McInnis, Carter C. Price, Susan C. Smith, and John S. Hollywood. *Predictive policing: the role of crime forecasting in law enforcement operations*. RAND Corporation, Santa Monica, CA, 2013.
16. Simone Silveti, Laura Nenzi, Ezio Bartocci, and Luca Bortolussi. Signal Convolution Logic. *ATVA*, 2018. arXiv: 1806.00238.